

Performance Enhancements

Extraction modules that operate solely on the minimized database D^1 – e.g. algebraic predicates – have negligible cost due to its minuscule size. However, modules that need to work with the original database D_I – e.g. NEP extraction – could incur significant overheads if D_I is large. Therefore, XPOSE employs the following techniques to reduce performance bottlenecks in these modules.

6.6 Correlated Sampling [17]

Correlated sampling is a technique that makes use of the schema join graph in the sampling process. This results in a higher probability of the sampled data satisfying the join predicates. It is used before the database minimization step to obtain a smaller D_I that produces a FIT-result, thereby reducing the iterations in the minimization. This is also useful in outer join queries, where random sampling is susceptible to producing tuples with mismatched keys.

6.7 View-based Database Minimization

We employ a minimization technique based on *virtual views*, which does not require copying the records of a table during the binary halving process. The views are created on the base table by utilizing system-generated tuple identifiers, which give the physical location of a row in the table – for instance, in PostgreSQL, this identifier is called `ctid` and consists of a block number and a record number within that block. The `ctid` of the first record of a table is $(0, 1)$. The number of records present in a block, n_b , is table-width dependent but computable from the schema. Based on n_b , we can estimate the `ctid` of the middle row of the table. For example, the following queries create a view containing roughly the upper half of table T :

```
Alter Table T Rename to Tdummy;
Create View T as
Select * From Tdummy
Where ctid between '(0,1)' and '(|Tdummy|/2nb,1)';
```

If a FIT-result is obtained, the view creation continues recursively with the upper half; if not, it shifts to a virtual view on the lower half. This reduction continues until a D^1 is achieved. The key improvement over the explicit halving approach is the avoidance of the time and space costs for materializing the intermediate tables.

6.8 Hash-Based Result Comparators

This component aims to efficiently verify the equality of the results of Q_H and Q_E over D_I . The direct but slow technique is to explicitly compute, using the EXCEPT command, the difference between the results in both directions – i.e. \mathcal{R}_H EXCEPT \mathcal{R}_E and \mathcal{R}_E EXCEPT \mathcal{R}_H , and verify that both are zero. This especially becomes a performance bottleneck in NEP extraction. Therefore, XPOSE employs the following hash-based result comparators instead. While PostgreSQL has several options for hash functions, we use the `HashText` hash function since it works across datatypes.

6.8.1 Global Hash Method: This method is used if the query has `ORDER BY` (\vec{O}) or `GROUP BY` attributes. It sorts \mathcal{R}_H and \mathcal{R}_E on remaining projection attributes wrt \vec{O} , then computes hashes on each result table. If they are equal, the result tables are the same.

6.8.2 Rolling Hash Method: This method is used if the query has no physical ordering. Here, we calculate a hash value for a relation

by applying a hash function to each tuple in the relation and then aggregating the results. In the PostgreSQL database, to get the rolling hash of the tuples present in the result set \mathcal{R}_H , we use:

```
Select SUM(rh_hashes.hashtext)
```

```
From (Select hashtext( $\mathcal{R}_H$ ::TEXT) From  $\mathcal{R}_H$ ) as rh_hashes;
```

The same evaluation is done for \mathcal{R}_E . Comparing the two hashes confirms or denies the unordered set equality of \mathcal{R}_H and \mathcal{R}_E .